

CS 2410 Graduate Computer Architecture

Fall 2016

Project 2: Dynamic Scheduling Simulator

Assigned: November 2, 2016.

Step 1 due: November 30, 2016 by 11:59 PM.

Step 2 due: December 8, 2016 by 1:00 PM.

1 Introduction

In this project, you will implement a simulation of dynamic instruction scheduling with Tomasulo's algorithm. Suppose we name the simulator, *tomsim*. Your simulator will mimic the behavior of a dynamic instruction scheduler to model the timing of the execution of instructions. The simulator does not need to model instruction semantics. *tomsim* will read an input program instruction trace and output statistics about the trace. You will integrate your simulator with OCCAM to setup, run and manage simulations.

This project will have two steps for the OCCAM part. First, you will need to complete the project using your own OCCAM instance (the VM). Second, you will incorporate your files into a demo OCCAM system that is online. Assuming you did a good job with the integration in your own VM, then it should be straightforward to incorporate the simulator into the demo system. You will use both your own VM and the demo system to demonstrate the project.

2 Overview

tomsim will input an instruction trace of X instructions. An instruction trace is simply an ordered list of the instructions executed by a program. The instruction trace for this project will not have any branch or jump instructions.

tomsim will also input a configuration that specifies the number of functional units (FU) of different types (integer, divider, multiplier, load and store), the number of reservation stations for each FU type, and the execution latency of each functional unit type.

Using the instruction trace and the configuration, *tomsim* will mimic the scheduling of Tomasulo's algorithm. Instructions will go through the Issue, Read Operand, Execute and Write Register steps of the algorithm.

Once all instructions are consumed from the trace, the simulator will output statistics.

3 Simulation Model

tomsim will need to model the pipeline stages for Tomasulo's algorithm (Issue, Read Operands, Execute and Write Register). The pipeline is single instruction in-order issue.

Issue: Instructions proceed in-order through Issue and Read Operands. If there is no reservation station available in Issue for an instruction, the pipeline stalls until a reservation station of the required type is available. If there is a reservation station available in Issue, then the reservation station is allocated to the instruction (marked busy).

Read Operand: After Issue, an instruction moves to Read Operand with an assigned reservation station. In Read Operand, the instruction reads any available source operands from the register file, renames its unavailable source registers, and renames its destination register (to the assigned reservation station).

If some operand is not available in Read Operand, the instruction waits until the operands are available in this stage. Instructions may bypass one another in Read Operand, depending on when their operands become available.

When an instruction becomes ready to execute (i.e., it has all of its operands), a check is done in Read Operand for the availability of a FU of the type needed to execute the instruction. If an appropriate functional unit is not available, then the instruction waits in Read Operand until a FU becomes available. Multiple instructions may be ready together and need the same FU type. If there are more ready instructions than available functional units of the required type, the oldest instructions are dispatched to the free FUs. Ties can be broken arbitrarily. The younger instructions continue to wait in Read Operand.

Execute: When operands are ready and a FU is available, an instruction is dispatched to the functional unit in Execute. A dispatched instruction stays in Execute, occupying the FU and reservation station, for the number of cycles specified by the FU's latency. For instance, if an Integer unit is specified to have latency 1, then an instruction spends 1 cycle in Execute after reading operands and being dispatched to the FU. In the best case, the integer instruction takes 4 cycles in total to go through the pipeline (Issue, Read Operands, 1-cycle Execute, Write Register). You do not need to model the semantics (the operation) of the instruction in Execute.

Write Register: Once an instruction finishes executing, it broadcasts the availability of its destination value to any waiting instructions and the register file. This broadcast happens in the Write Register stage. It takes 1 clock cycle. Thus, a consumer instruction waiting on a register value can potentially be dispatched to execute on the cycle after the broadcast. You may assume an infinite number of common data busses (CDB)! That is, you do not need to model the potential structural hazard for a CDB. Note that multiple instructions can proceed through Write Register on the same clock cycle. In Write Register, the functional unit and reservation station occupied by an instruction are released. You do not need to model the actual data values, but you will need to model the renaming associated with the completion of an instruction.

Pipeline Model: There are five functional unit types: Integer, Multiplier, Divider, Load, and Store. The table below gives the opcode types that are executed by each functional unit type.

Type	Opcodes
Integer	add, sub, nor, and, lis, liz, lui, put, halt
Divider	div, exp, mod
Multiplier	mul
Load	lw
Store	sw

Each functional unit has reservation stations associated with it. The reservation stations are associated by functional unit type. For instance, the reservation stations for an Integer functional unit are separate from the reservation stations for a Divider functional unit. However, all functional units of the same type share the same set of reservation stations. Thus, an instruction can be dispatched from the reservation station of some type to any functional units of that type. For instance, a pipeline might have 2 Integer units sharing 8 reservation stations. The organization for this case is shown in Figure 1. An instruction from the 8 reservation stations can be dispatched to either of the Integer functional units.

4 Project Requirements

The simulator should be runnable from the command line. It will take three arguments:




Figure 1: 8 reservation stations with two Integer functional units.

```
$tomsim trace.t configuration.json output.json
```

In this example, `trace.t` is the input trace, `configuration.json` is the configuration of the pipeline, and `output.json` is the output result.

4.1 Input Trace

The input trace is a list of ***X*** encoded instructions. It has the same format and encodings as project 1. You can re-use your input program code from project 1. However, the input trace will not have any branches or jump instructions.

For testing, you can use your ***X*** simulator. Simply add output statements to your simulator that outputs the instructions executed to a file (one instruction in hex format per line). When a branch or jump is executed, do not output anything. This output file is now an instruction trace that can be input to ***tomsim***.

Here is an example trace:

```
800A
7000
89FF
7020
9100
7020
6800
```

4.2 Configuration

The simulator will be configured with a JSON file. The configuration will permit specifying the number of functional units of different types, the number of reservation stations associated with a functional unit type, and the latency of executing in the functional unit. The JSON has the format:

```
{"integer":[{"number":2,"resnumber":4,"latency":1}],
"divider":[{"number":1,"resnumber":2,"latency":20}],
```

```
"multiplier":[{"number":2,"resnumber":4,"latency":10}],
"load":[{"number":1,"resnumber":8,"latency":100}],
"store":[{"number":1,"resnumber":4,"latency":2}]}
```

This configuration specifies the following:

1. There are 2 Integer FUs and 4 reservation stations. An integer operation takes 1 clock cycle.
2. There is 1 Divider FU unit with 2 reservation stations. A divide executes in 20 clock cycles.
3. There are 2 Multiplier FUs with 4 reservation stations. A multiply takes 10 clock cycles.
4. There is 1 Load FU. This unit has 8 reservation stations and takes 100 clock cycles to read memory.
5. There is 1 Store FU, which has 4 reservation stations. It takes 2 clock cycles to write memory.

4.3 Output

The simulator will output statistics in JSON about execution. The output statistics are: number of clock cycles for the whole trace, the number of instructions executed by each FU, the number of stall cycles due to structural hazards, and the number of operands read from the register file. The output JSON has the format:

```
{"cycles" : 100,
"integer" : [{ "id" : 0, "instructions" : 10 }, { "id" : 1, "instructions" : 8 }],
"multiplier" : [{ "id" : 0, "instructions" : 1 }, { "id" : 1, "instructions" : 0 }],
"divider" : [{ "id" : 0, "instructions" : 2 }],
"load" : [{ "id" : 0, "instructions" : 3 }],
"store" : [{ "id" : 0, "instructions" : 3 }],
"reg reads" : 9,
"stalls" : 13 }
```

In this output, the field `id` distinguishes multiple functional units of the same type. The field `instructions` is the number of instructions. The entries `reg reads`, `stalls` and `cycles` report the number of register file reads, the number of pipeline structural hazard stalls and the number of cycles for the simulation.

4.4 Using OCCAM

There will be two steps for integrating with OCCAM.

Step 1: You will integrate your simulator into your own OCCAM instance (you can use the same VM that you used previously). Similar to the first project, the integration will involve writing “wrappers” for the configuration input, building the simulator, running the simulator, and parsing the statistics output. The OCCAM integration should allow configuring the simulation (changing the instruction latencies), selecting different input programs, running the program, and viewing the output statistics. You can probably reuse your project 1 to simplify the integration.

Step 2: Once you have integrated your simulator into your own OCCAM instance, you will incorporate it into a live online demo system. This step should go smoothly, as long as you have done a good job of integrating in your own instance and well tested the simulator and the configuration. More details about this step (including a brief “howto” guide) will be made available closer to the due date.

The project demo will be done after Step 2 is completed. The demo will show both your integration in your own OCCAM instance as well as the online system.

Note: Step 1 and Step 2 have different due dates. See below.

4.5 Programming Language, Libraries, etc

You are welcome to use any programming language that you wish. However, the simulator must be runnable from within OCCAM. Nevertheless, I suggest that you use C++ (or C). We selected JSON for the configuration and output files for two reasons. First, using JSON for these files will make it much easier to integrate your simulator with OCCAM. Second, there are packages (e.g., `jsoncpp`, `rapidjson`) available that you can use to parse JSON. You are welcome to use these libraries to simplify the effort to read/write the files.

5 Turning in the Project

For grading, you will both turn in the project by email (as explained below) and schedule a demo of the project with Bruce Childers. Your project will be graded primarily during the demo with possibly additional follow-up after the demo.

5.1 What to Submit

Prior to scheduling the demo, you must submit all files for the project. You must submit the files:

- `tommy.cc` (the source file)
- all OCCAM related files
- `README.txt` (a help file; see below)

These files should be put into a single archive; you may use either a gzipped tar-ball, or a zip file. Please name your submission with your last name. For instance, `childers-project2.zip` is a good file name for Bruce Childers' submission of project 2. I will use a Mac to unpack your submission. If I cannot unpack the submission, or there are any missing files, I will not grade the project and you may receive a 0.

In the README, put your name and e-mail address at the top of the file. The README file should list what works (i.e., instructions implemented) and any known problems. If you have known problems/issues (bugs!), then you should clearly specify the problems in this file. The explanation and bug list are critical to grading. So, if you're uncertain whether to include something, then err on the side of writing too much and just include it.

Step 1 is due: November 30, 2016 by 11:59 PM. Submit your files by email.

Step 2 is due: December 8, 2016 by 1:00 PM. Demos will be held on this date.

5.2 Where to Submit

Submit your files to Bruce Childers, email `childers@cs.pitt.edu`).

5.3 Demonstration

After Step 2, you will demo your simulator for grading. I will give you some test files (configurations and traces) to run in the demonstration. These files will *not* be provided prior to the demonstration. I will give them to you on a USB key. You should be able to load the test files into your OCCAM instance (running your simulator), configure the simulator from the OCCAM interface, and run the simulation. You should be comfortable with demonstrating the project in OCCAM, such as configuring the simulator with OCCAM, viewing results in OCCAM, and creating graphs from the results (e.g., plot a histogram of the instruction counts).

For the online demo, you will configure a “work flow” of a trace generator that executes a **X** program and outputs a trace, which is input to your simulator.

All projects will be demonstrated on December 8, 2016. Each demo will be 30 minutes. A sign-up sheet with appointment times will be made available prior to the demo date. If you have a

class conflict with the demo day/times, we can make accommodations. If I have not received your file submission by the due day/time, I will not grade the project and you may receive a 0.

Because the demos are intended to be short, please be prepared to quickly and smoothly show what you have done! The time allotted for each demo is strict.

6 Collaboration

In accordance with the policy on individual work, this project is strictly an individual effort. You must not collaborate with a partner. It is your responsibility to secure and back up your files.

6.1 Grading the Project

I will grade the project with multiple configurations and traces covering many (if not all) of X instructions during the demo. Your ability to easily and smoothly demonstrate the project in OCCAM will also be part of the evaluation. Primarily, the concern is functional correctness and implementation of the project requirements. To ensure that your processor is working as expected, you will find it very useful to write your own test case programs to cover all instructions.

7 Hints

The project needs a way to manage “simulation time”. The simulation time is simply a count of cycles. The simulation proceeds in steps, which increments the simulation clock cycle count by one.

On each simulated clock cycle, multiple actions need to be performed to model the movement and execution of instructions in the pipeline. The actions need to happen according to pipeline conditions (e.g., operand availability) and latency of the actions. With this in mind, you will need some way to cause simulation actions to happen on specific simulation clock cycles.

While there are many ways of managing the simulation time and actions, a classic approach is discrete event simulation. In this method, a list of events with timestamps is maintained. The simulation proceeds in steps of a single clock cycle. On each clock cycle, the list is checked for events that have a timestamp matching the current time. A matching event is removed from the list and handled on the current clock cycle.

It will be useful to allow multiple event types in the list, where each event type is a different action in the life of an instruction. An instruction will cause some event to be in the list, which will be processed on a matching clock cycle to update the movement of the instruction.

As an example, consider an Instruction in Read Operands. This instruction might have put a Read Operand event in the event list during Issue. The handler for the Read Operand event would read the register file (when the simulation clock cycle matches the event’s timestamp). If all operands can be read from the register file and a FU is available, a new Execute event would be inserted (i.e., scheduled) for the next clock cycle. If the operands are not all ready or there is no FU, then a new event would be scheduled to Read Operands again on the next cycle. In this way, the instruction remains in Read Operands until the conditions to execute are satisfied. Other actions for in-flight instructions can be handled similarly. You will probably need at least 4 event types: Issue, Read Operands, Execute, and Write Result. Note, you may wish to have more event types, such as a Wait event type when an instruction is waiting on operands/FU.

Most events can be inserted with a timestamp corresponding to the next clock cycle. An Execute event, however, might be inserted with a further-away timestamp that corresponds to the latency of the instruction (i.e., current time + latency of instruction).

You may also want to keep the event list ordered by timestamp (soonest first). This can help the efficiency of processing the event list since you can terminate processing the list on a simulation cycle when a timestamp is found that is greater than the current simulation clock cycle.

You may also find it useful to process events in order by their type. You may wish to process

events for later instructions in the pipeline before events of earlier instructions. For example, it could be useful to process Write Register before Read Operand.

I strongly recommend that you plan your simulation before trying to implement it. The data structures that you choose will likely have a big influence on your required effort. So, think carefully here!! I also strongly recommend that you implement and test in steps. It is very difficult to debug problems in this type of simulation; proceeding in small steps and carefully testing your implementation in an incremental fashion can be a big life saver! Lastly, you will probably want to write carefully crafted test cases. You can craft these test cases to exhibit certain timing behavior, which can then be checked to test your simulator.